

# Neural Tree Expansion for Prioritized Multi-Agent Path Deconfliction

Khushal Brahmbhatt  
Andrew Festa  
Kavinda Senewiratne

Kavi - outline of slides

Problem:

- picture of wildland fire

- Consequences of wildfires - lives lost, assets destroyed, getting worse every year (global warming)

- why monitoring fronts is important (high risk decision making - minimal error, latest info needed)

- Current methods and why they are insufficient (satellite - low temporal res, terrestrial - no communication)

Related Work (method and issues with them)

- Heuristic based methods - prioritization heuristics (path known before hand, priority assigned on capability of robot and not the task)

- Conflict based search - fail to consider priority, centralized

- Mixed integer solutions

Solution

- NTE - imitation learning

- Pictorial algo

- Explain each component

- Briefly touch on what we are changing in NTE and why - path confliction

Contributions

- Reward function

Try get it down to 3:30

# Background - Wildland Fires




- Countless civilian and firefighter **lives lost**.
- \$10Bn in assets damaged in the US.
- Destruction of natural habitats.

# Problem - Monitoring Firefronts with Multiple Robots



- Fire Fronts help predict spread of fire.
- Decision making under extreme conditions.
- **Paths could conflict.**

# Path Conflict



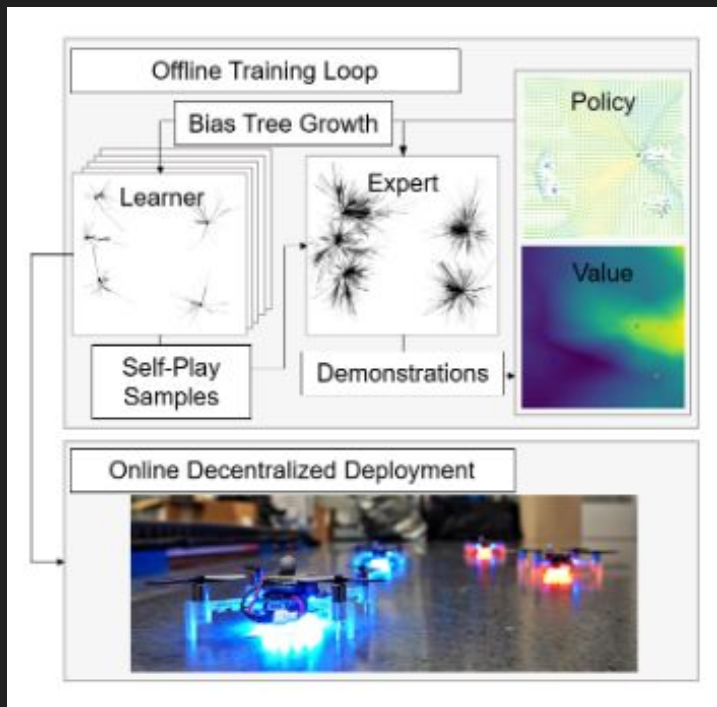
## Our Focus

1. Decentralization
2. Focus on priority

# Related Work (Multi-Agent Path Deconfliction)

- Heuristic based methods
  - Path known beforehand
  - Priority assigned on capability of robot and not the task
- Conflict based search
  - Fail to consider priority
  - Centralized Mixed integer solutions

# Method

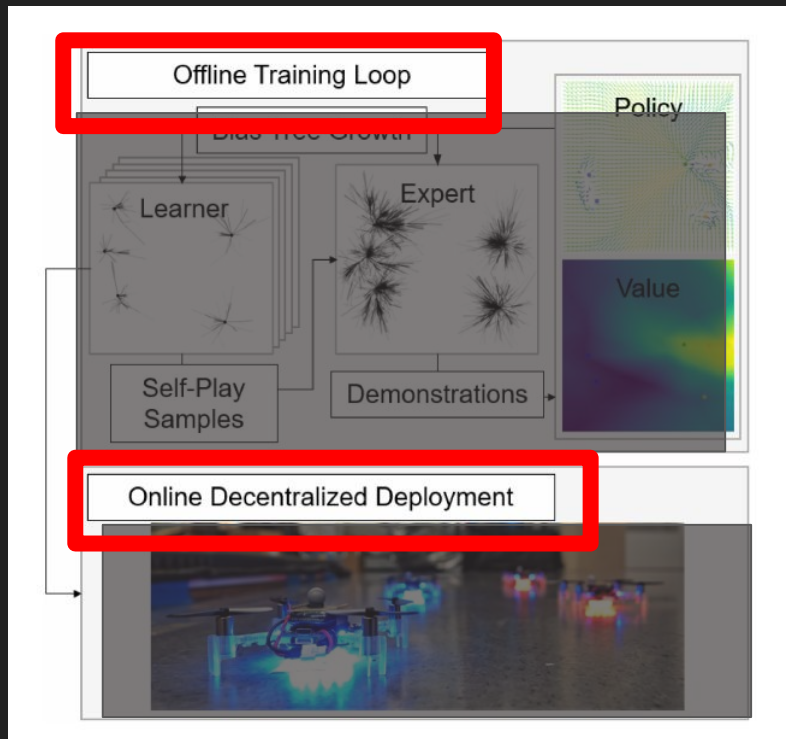


## Neural Tree Expansion for Multi-Robot Planning in Non-Cooperative Environment

Benjamin Riviere, Wolfgang Hoenig, Matthew Anderson, and Soon-Jo Chung

(Imitation Learning)

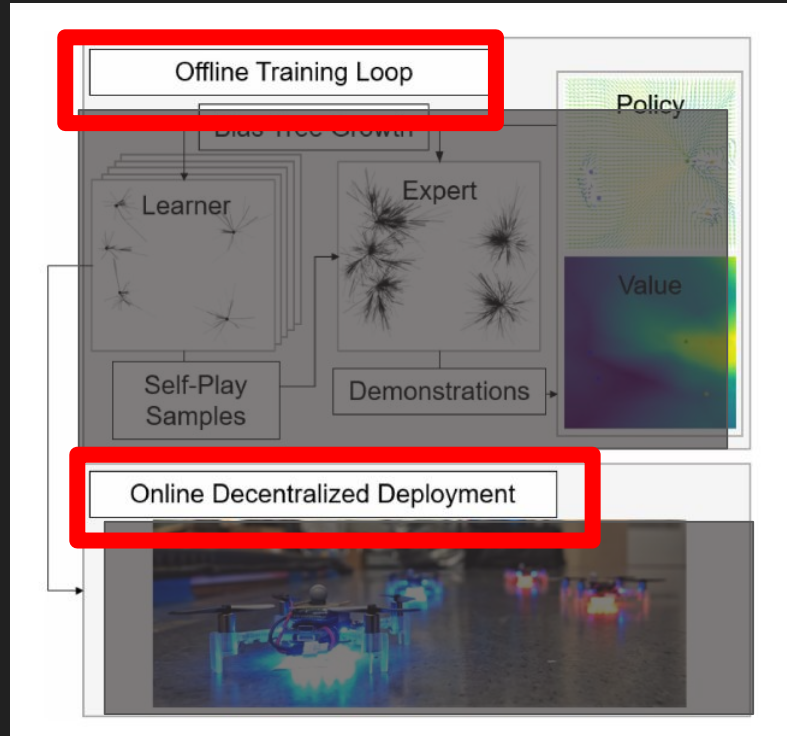
# Method



Policies  
for  
Robots

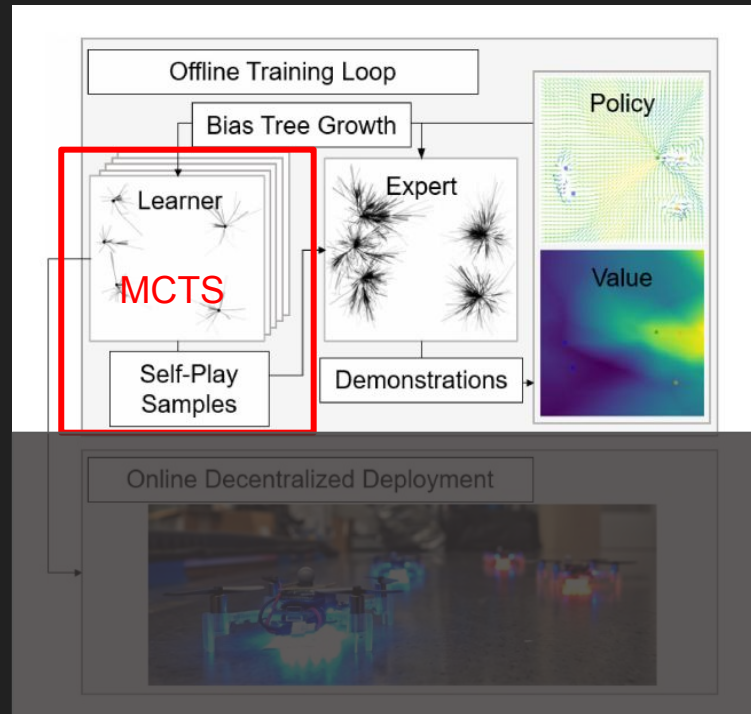


# Method



Source: Neural Tree Expansion for Multi-Robot Planning in Non-Cooperative Environment

# Method



Source: Neural Tree Expansion for Multi-Robot Planning in Non-Cooperative Environment

# Method



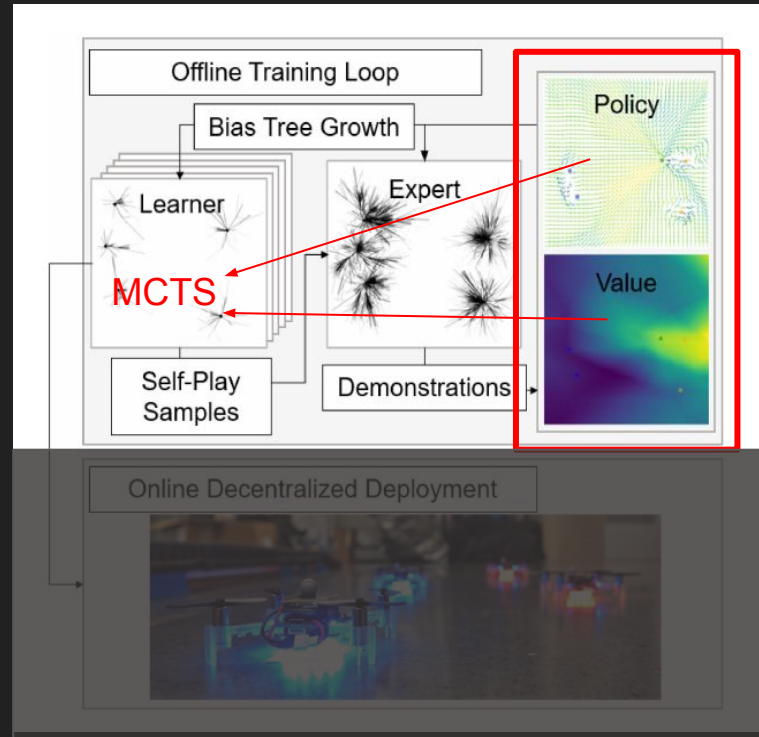
Source: Neural Tree Expansion for Multi-Robot Planning in Non-Cooperative Environment

# Method



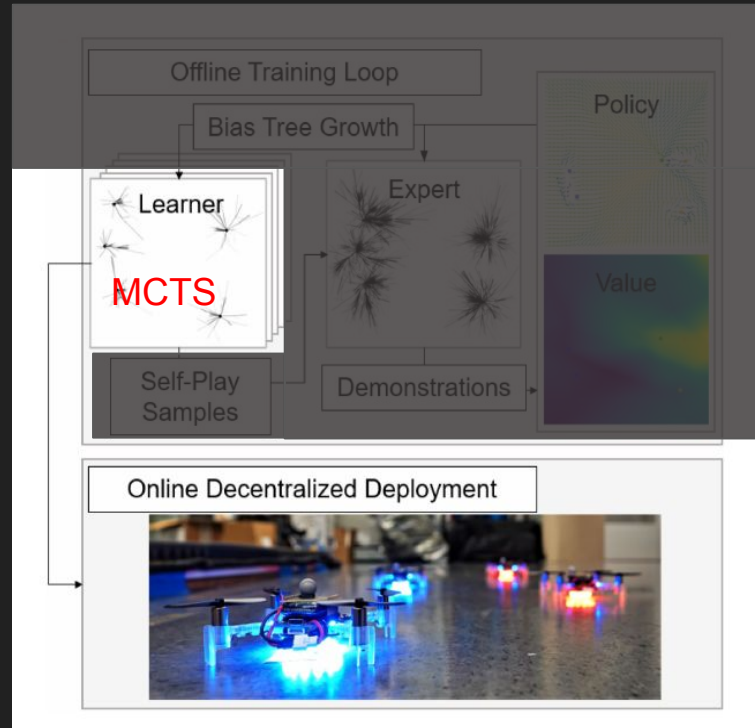
Source: Neural Tree Expansion for Multi-Robot Planning in Non-Cooperative Environment

# Method



Source: Neural Tree Expansion for Multi-Robot Planning in Non-Cooperative Environment

# Method



Source: Neural Tree Expansion for Multi-Robot Planning in Non-Cooperative Environment

# Our Contribution - Reward Function



# Approach

- PUCT search
- Agents trained with local knowledge (sensing radius)
- Polar representation for surroundings and movement
- Explicit agent prioritization in reward function



# Polynomial Upper Confidence Trees (PUCT)


- Used in AlphaZero
- Variant of MCTS
  - Selection
  - Expansion
  - Simulation
  - Backpropagation
- Think MiniMax

# State Representation

$$s_t = \langle x_{ri}, y_{ri}, p_{ri}, x_{gi}, y_{gi}, [d_1 \dots d_8], [p_1 \dots p_8] \rangle$$

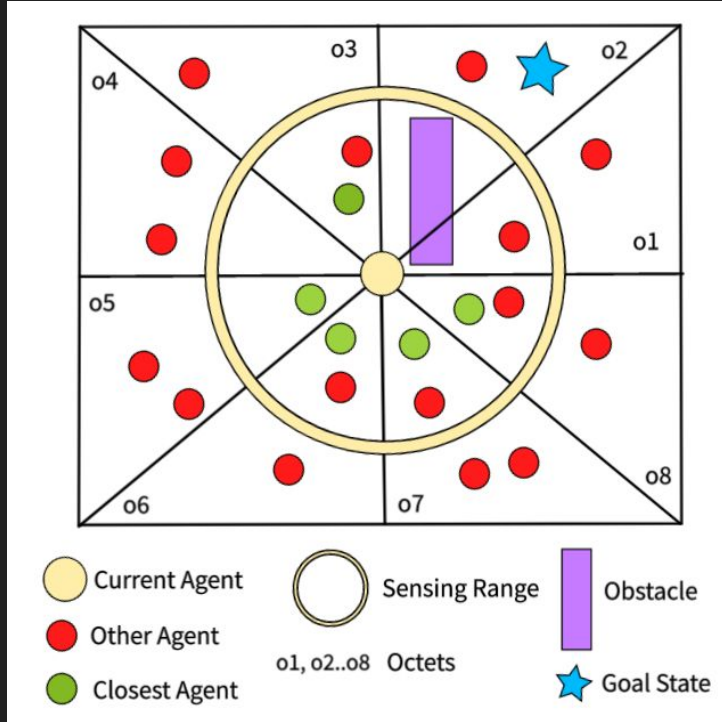
- $x_{ri}$  x-coordinate of robot  $i$
- $y_{ri}$  y-coordinate of robot  $i$
- $p_{ri}$  priority of robot  $i$
- $x_{gi}$  x-coordinate of the goal of robot  $i$
- $y_{gi}$  y-coordinate of the goal of robot  $i$
- $[d_1 \dots d_8]$  distance to the closest agent in each octant  $[o_1 \dots o_8]$
- $[p_1 \dots p_8]$  priority of the closest agent in each octant  $[o_1 \dots o_8]$

# State Examples




- $x_{ri}$  0
- $y_{ri}$  0
- $p_{ri}$  1
- $x_{gi}$  6
- $y_{gi}$  1
- $[d_1 \dots d_8]$  [2.8, 4, 1.3, 4, 1, 0.9, 1.1, 1.2]
- $[p_1 \dots p_8]$  [1, 0, 0.4, 0, 0.2, 0.6, 0.5, 0.5]

# State Examples



- $x_{ri}$  0
- $y_{ri}$  0
- $p_{ri}$  1
- $x_{gi}$  8
- $y_{gi}$  1
- $[d_1 \dots d_8]$  [0.3, 0.7, 1.3, 4, 1, 0.9, 1.1, 1.2]
- $[p_1 \dots p_8]$  [2, 2, 0.4, 0, 0.2, 0.6, 0.5, 0.5]

# State Examples



- Information in each octant considers closest agent in that octant
- Regions with no agents?
- Regions with an obstacle?

# Reward Function


$$R_x = \frac{1}{d_g} + \sum_{i=1}^8 (P_x - P_i) \left( 1 - \frac{d_{xi}}{r_x} \right)$$

- $R_x$  reward for agent  $x$
- $d_g$  distance between agent  $x$  and its goal
- $P_x$  priority for agent  $x$
- $P_i$  closest agent in octant  $i$
- $d_{xi}$  distance between agent  $x$  and the closest agent in octant  $i$
- $r_x$  sensing range of agent  $x$

## But that's just math

- Incentivize moving towards goal
- Yield right of way when approaching another agent with a higher priority
- Continue with current trajectory if approaching agent with lower priority
- Avoid obstacles


# Reward Examples



•	$a_0$	1
•	$a_1$	1
•	$a_2$	1
•	$a_3$	1
•	$a_4$	1
•	$a_5$	1
•	$a_6$	1
•	$a_7$	1
•	$a_8$	1




# Reward Examples




•	$a_0$	1
•	$a_1$	1
•	$a_2$	1
•	$a_3$	1
•	$a_4$	1
•	$a_5$	1
•	$a_6$	1
•	$a_7$	1
•	$a_8$	1

# Reward Examples



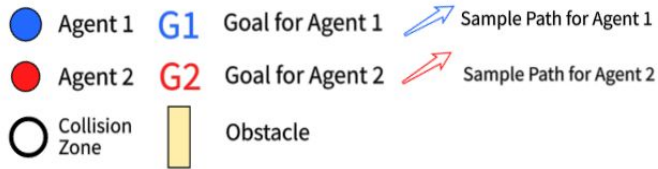

- Move towards goal
- Move in direction with no agents
- Move away from agents with high priority
- Avoid obstacles!

# Experiments





- Training iterations: 10/50
  - Expert demos epochs: 100
  - Expert demos iterations: 100
  - Search depth : 25
- Learning rate: 0.001

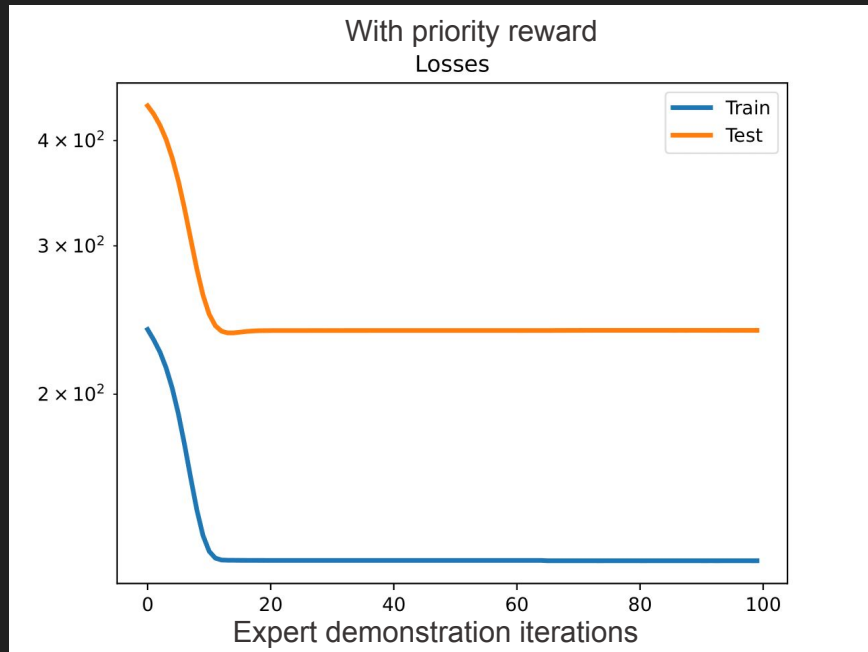
# Experiments



# Results - Trajectories



# Results - Value Learning curves



# Current Limitations and Next Steps

- Priority weightage
- All agents have the same speed
- All agents have the same sensing range

# More! More! More!

- Scale up number of agents
- Larger map with obstacles
- System throughput?
- Common goals?
- Variable speeds/Continuous action space?



Questions?